

SKLearn using python

Jiarui Groves

4/8/2023

CS 4375: Machine Learning

```
from google.colab import files
```

```
uploaded = files.upload()
```

[Browse...](#) Auto.csv

Auto.csv(application/vnd.ms-excel) - 17859 bytes, last modified: n/a - 100% done

Saving Auto.csv to Auto (1).csv

```
import io
import pandas as pd
df = pd.read_csv(io.BytesIO(uploaded['Auto.csv']))
df.head()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	origin	r
0	18.0	8	307.0	130	3504	12.0	70.0	1 chev	chev
1	15.0	8	350.0	165	3693	11.5	70.0	1 b	sky
2	18.0	8	318.0	150	3436	11.0	70.0	1 plym	sat

```
#df['Length'].value_counts
print(len(df))
print('\nDimensions of data frame:', df.shape)
```

392

Dimensions of data frame: (392, 9)

Describe on auto data shows

✓ 4s completed at 2:25 PM



3. Range of year is from 70 to 82 years old

```
df.describe()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	
count	392.000000	392.000000	392.000000	392.000000	392.000000	391.000000	392
mean	23.445918	5.471939	194.411990	104.469388	2977.584184	15.554220	70
std	7.805007	1.705783	104.644004	38.491160	849.402560	2.750548	7
min	9.000000	3.000000	68.000000	46.000000	1613.000000	8.000000	70
25%	17.000000	4.000000	105.000000	75.000000	2225.250000	13.800000	70
50%	22.750000	4.000000	151.000000	93.500000	2803.500000	15.500000	70
75%	29.000000	8.000000	275.750000	126.000000	3614.750000	17.050000	70
max	46.600000	8.000000	455.000000	230.000000	5140.000000	24.800000	80

```
df.dtypes
```

```

mpg          float64
cylinders     int64
displacement  float64
horsepower    int64
weight        int64
acceleration  float64
year          float64
origin        int64
name          object
dtype: object

```

Type of data contained are summarized above.

```

# Change cylinder to categorial type
df.cylinders=df.cylinders.astype('category').cat.codes
print("\n Cylinders after category conversion is now \n",df.cylinders)
df.origin = df.origin.astype('category')
print("\n Origin after category conversion is now \n", df.origin)
# Now the new dtype on data
print(" Here are the data after conversion\n:", df.dtypes)
print(" Converted data: \n",df.head)

```

Cylinders after category conversion is now

```
0      4
1      4
2      4
3      4
4      4
```

```
..
```

```
387    1
388    1
389    1
390    1
391    1
```

Name: cylinders, Length: 392, dtype: int8

Origin after category conversion is now

```
0      1
1      1
2      1
3      1
4      1
```

```
..
```

```
387    1
388    2
389    1
390    1
391    1
```

Name: origin, Length: 392, dtype: category

Categories (3, int64): [1, 2, 3]

Here are the data after conversion

```
: mpg          float64
cylinders      int8
displacement   float64
horsepower     int64
weight         int64
acceleration    float64
year           float64
origin         category
name           object
```

dtype: object

Converted data:

	<bound method NDFrame.head of	mpg	cylinders	displacement	horsepower	weight
0	18.0 4 307.0	130	3504	12.0	70.0	
1	15.0 4 350.0	165	3693	11.5	70.0	
2	18.0 4 318.0	150	3436	11.0	70.0	
3	16.0 4 304.0	150	3433	12.0	70.0	
4	17.0 4 302.0	140	3449	NaN	70.0	
..	
387	27.0 1 140.0	86	2790	15.6	82.0	
388	44.0 1 97.0	52	2130	24.6	82.0	
389	32.0 1 135.0	84	2295	11.6	82.0	
390	28.0 1 120.0	79	2625	18.6	82.0	
391	31.0 1 119.0	82	2720	19.4	82.0	

```
origin          name
0      1  chevrolet chevelle malibu
```

```

1      1      buick skylark 320
2      1      plymouth satellite

```

```

df.isnull().sum()
# Original data oconsists 3 N/As
# Drop NAs in dataset
df = df.dropna()
df.shape

```

```

(389, 9)

```

New dimension of data decreased from 392 to 389 after the removal of 3 NAs

```

#mpg_high = df[['mpg']]
#mpg_high
mean = df["mpg"].mean()
mean    #find mean of column mpg

```

```

23.445918367346938

```

Printed entire column of mpg_high after making it.

```

df.loc[df['mpg'] > mean, 'mpg_high'] = '1'
df.loc[df['mpg'] <= mean, 'mpg_high'] = '0'
#print(df.loc[:'mpg'])
print(df.mpg_high)
print(df[['mpg_high']].to_string(index=False))

#print(df.mpg)

0      0
1      0
2      0
3      0
6      0
..
387    1
388    1
389    1
390    1
391    1
Name: mpg_high, Length: 389, dtype: object
mpg_high
0
0
0
0
0

```

```
#Data after dropping original mpg is printed out
df.drop('mpg',inplace = True, axis = 1)
print(df.head)
```

4/8/2023, 2:28 PM


388	1	97.0	52	2130	24.6	82.0	2
389	1	135.0	84	2295	11.6	82.0	1
390	1	120.0	79	2625	18.6	82.0	1
391	1	119.0	82	2720	19.4	82.0	1

	name	mpg_high
0	chevrolet chevelle malibu	0
1	buick skylark 320	0
2	plymouth satellite	0
3	amc rebel sst	0
4	ford torino	0
..
387	ford mustang gl	1
388	vw pickup	1
389	dodge rampage	1
390	ford ranger	1
391	chevy s-10	1

[392 rows x 9 columns]>

```
import seaborn as sb
from sklearn import datasets
sb.catplot(x="mpg_high", kind="count", data=df)
```

<seaborn.axisgrid.FacetGrid at 0x7f99bf399940>




```
# seaborn relplot with horsepower on the x axis, weight on the y axis, setting hue or style
```

```
#Relaional plot shoes that there are a esignificant trend between weight and horsepower dat
```

```
sb.relplot(x="horsepower", y="weight", data=df, hue=df.mpg_high, style=df.mpg_high)
```


```
<seaborn.axisgrid.FacetGrid at 0x7f99ba93fb80>
```




```
#seaborn boxplot with mpg_high on the x axis and weight on the y axis
```

```
sb.boxplot(x = "mpg_high", y="weight", data=df)
```

```
<Axes: xlabel='mpg_high', ylabel='weight'>
```





Data tells us that among higher MPG/ more efficient vehicles, there are more higher weight ones while in vehicle with lower mpg/less efficient ones, there are no significant trends in weight

Created train and test of 80/20 percent. Obtained of train size of 313 rows by 8 columns, and test size of 79 rows by 8 columns

```
from sklearn.model_selection import train_test_split
X = df.loc[:, ['cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'year', 'origin']]
y = df.mpg_high
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
```

```
print('train size:', X_train.shape)
print('test size:', X_test.shape)
```

```
train size: (311, 7)
test size: (78, 7)
```

Logistic regression train a logistic regression model using solver lbfgs b. test and evaluate c. print metrics using the classification report

```
from sklearn.linear_model import LogisticRegression
#convert into categorical data
clf = LogisticRegression()
clf.fit(X_train, y_train)
clf.score(X_train, y_train)
```

```
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning:
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```


Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
0.9035369774919614
```

Logistic regression shows around 90% accuracy

Copied df to a new dataset newdf to perform decision tree classification without disturbing original dataset Using cylinders, year and horsepower as predictors

```
newdf = df.copy()
print(newdf.dtypes)
newdf.cylinders = newdf.cylinders.astype('category').cat.codes
newdf.year = newdf.year.astype('category').cat.codes
newdf.horsepower = newdf.horsepower.astype('category').cat.codes
print(newdf.dtypes)
```

```
mpg          float64
cylinders    int64
displacement float64
horsepower   int64
weight       int64
acceleration float64
year         float64
origin       int64
name         object
mpg_high     object
dtype: object
mpg          float64
cylinders    int8
displacement float64
horsepower   int8
weight       int64
acceleration float64
year         int8
origin       int64
name         object
mpg_high     object
dtype: object
```

#Decision tree using the same train and test dataset

```
X = newdf.loc[:, ['cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'year', 'mpg_high']]
y = newdf.mpg_high
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
```

```
print('train size:', X_train.shape)
```

```
print('train size: ', X_train.shape)
print('test size: ', X_test.shape)
```

```
train size: (311, 7)
test size: (78, 7)
```

```
from sklearn.tree import DecisionTreeClassifier
```

```
clf = DecisionTreeClassifier()
clf.fit(X_train, y_train)
```

```
▼ DecisionTreeClassifier
DecisionTreeClassifier()
```

```
# make predictions
```

```
pred = clf.predict(X_test)
```

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
print('accuracy score: ', accuracy_score(y_test, pred))
#print('precision score: ', precision_score(y_test, pred))
#print('recall score: ', recall_score(y_test, pred))
#print('f1 score: ', f1_score(y_test, pred))
```

```
accuracy score: 0.8974358974358975
```

```
from sklearn.metrics import confusion_matrix
# COnfusion matrix of decision tree gives 4 false positive, and 4 false negative
confusion_matrix(y_test, pred)
```

```
array([[39, 5],
       [ 3, 31]])
```

```
# Report shows result is slightly worse then the Logistic regtression model. with 3 false r
```

```
from sklearn.metrics import classification_report
print(classification_report(y_test, pred))
```

```

              precision    recall  f1-score   support

0               0.92      0.90      0.91         44
1               0.89      0.81      0.85         34
```

	0	0.93	0.89	0.91	44
	1	0.86	0.91	0.89	34
accuracy				0.90	78
macro avg		0.89	0.90	0.90	78
weighted avg		0.90	0.90	0.90	78

```
from sklearn import tree
tree.plot_tree(clf)
```

```
[Text(0.71328125, 0.9444444444444444, 'x[1] <= 190.5\ngini = 0.5\nsamples =
311\nvalue = [159, 152]'),
Text(0.5015625, 0.8333333333333334, 'x[2] <= 96.5\ngini = 0.292\nsamples =
180\nvalue = [32, 148]'),
Text(0.228125, 0.7222222222222222, 'x[1] <= 113.5\ngini = 0.188\nsamples =
152\nvalue = [16, 136]'),
Text(0.05, 0.6111111111111112, 'x[0] <= 3.5\ngini = 0.041\nsamples = 96\nvalue =
[2, 94]'),
Text(0.025, 0.5, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(0.075, 0.5, 'x[2] <= 87.5\ngini = 0.021\nsamples = 95\nvalue = [1, 94]'),
Text(0.05, 0.3888888888888889, 'gini = 0.0\nsamples = 82\nvalue = [0, 82]'),
Text(0.1, 0.3888888888888889, 'x[4] <= 18.8\ngini = 0.142\nsamples = 13\nvalue =
[1, 12]'),
Text(0.075, 0.2777777777777778, 'gini = 0.0\nsamples = 11\nvalue = [0, 11]'),
Text(0.125, 0.2777777777777778, 'x[4] <= 19.3\ngini = 0.5\nsamples = 2\nvalue = [1,
1]'),
Text(0.1, 0.16666666666666666, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(0.15, 0.16666666666666666, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(0.40625, 0.6111111111111112, 'x[5] <= 75.5\ngini = 0.375\nsamples = 56\nvalue
= [14, 42]'),
Text(0.275, 0.5, 'x[4] <= 16.75\ngini = 0.475\nsamples = 18\nvalue = [11, 7]'),
Text(0.2, 0.3888888888888889, 'x[1] <= 115.5\ngini = 0.444\nsamples = 9\nvalue =
[3, 6]'),
Text(0.175, 0.2777777777777778, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(0.225, 0.2777777777777778, 'x[4] <= 16.0\ngini = 0.245\nsamples = 7\nvalue =
[1, 6]'),
Text(0.2, 0.16666666666666666, 'gini = 0.0\nsamples = 5\nvalue = [0, 5]'),
Text(0.25, 0.16666666666666666, 'x[3] <= 2338.5\ngini = 0.5\nsamples = 2\nvalue =
[1, 1]'),
Text(0.225, 0.05555555555555555, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(0.275, 0.05555555555555555, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(0.35, 0.3888888888888889, 'x[4] <= 17.5\ngini = 0.198\nsamples = 9\nvalue =
[8, 1]'),
Text(0.325, 0.2777777777777778, 'x[2] <= 79.0\ngini = 0.444\nsamples = 3\nvalue =
[2, 1]'),
Text(0.3, 0.16666666666666666, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(0.35, 0.16666666666666666, 'gini = 0.0\nsamples = 2\nvalue = [2, 0]'),
Text(0.375, 0.2777777777777778, 'gini = 0.0\nsamples = 6\nvalue = [6, 0]'),
Text(0.5375, 0.5, 'x[4] <= 21.85\ngini = 0.145\nsamples = 38\nvalue = [3, 35]'),
Text(0.475, 0.3888888888888889, 'x[2] <= 93.5\ngini = 0.105\nsamples = 36\nvalue =
[2, 34]'),
Text(0.425, 0.2777777777777778, 'x[3] <= 2880.0\ngini = 0.059\nsamples = 33\nvalue
= [1, 32]'),
```

```

Text(0.4, 0.16666666666666666, 'gini = 0.0\nsamples = 26\nvalue = [0, 26]'),
Text(0.45, 0.16666666666666666, 'x[3] <= 2920.0\ngini = 0.245\nsamples = 7\nvalue =
[1, 6]'),
Text(0.425, 0.05555555555555555, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(0.475, 0.05555555555555555, 'gini = 0.0\nsamples = 6\nvalue = [0, 6]'),
Text(0.525, 0.27777777777777778, 'x[4] <= 14.5\ngini = 0.444\nsamples = 3\nvalue =
[1, 2]'),
Text(0.5, 0.16666666666666666, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(0.55, 0.16666666666666666, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(0.6, 0.38888888888888889, 'x[5] <= 77.5\ngini = 0.5\nsamples = 2\nvalue = [1,
1]'),
Text(0.575, 0.27777777777777778, 'gini = 0.0\nsamples = 1\nvalue = [1, 0]'),
Text(0.625, 0.27777777777777778, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(0.775, 0.7222222222222222, 'x[5] <= 78.5\ngini = 0.49\nsamples = 28\nvalue =
[16, 12]'),
Text(0.75, 0.6111111111111112, 'x[3] <= 2702.5\ngini = 0.266\nsamples = 19\nvalue =
[16, 3]'),
Text(0.725, 0.5, 'x[5] <= 77.5\ngini = 0.5\nsamples = 6\nvalue = [3, 3]'),
Text(0.7, 0.38888888888888889, 'x[3] <= 2630.0\ngini = 0.375\nsamples = 4\nvalue =
[3, 1]'),
Text(0.675, 0.27777777777777778, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]'),
Text(0.725, 0.27777777777777778, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(0.75, 0.38888888888888889, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(0.775, 0.5, 'gini = 0.0\nsamples = 13\nvalue = [13, 0]'),
Text(0.8, 0.6111111111111112, 'gini = 0.0\nsamples = 9\nvalue = [0, 9]'),
Text(0.925, 0.8333333333333334, 'x[4] <= 21.6\ngini = 0.059\nsamples = 131\nvalue =
[127, 4]'),
Text(0.9, 0.7222222222222222, 'x[5] <= 80.5\ngini = 0.045\nsamples = 130\nvalue =
[127, 3]'),
Text(0.85, 0.6111111111111112, 'x[2] <= 83.0\ngini = 0.016\nsamples = 124\nvalue =
[123, 1]'),
Text(0.825, 0.5, 'x[2] <= 79.5\ngini = 0.375\nsamples = 4\nvalue = [3, 1]'),
Text(0.8, 0.38888888888888889, 'gini = 0.0\nsamples = 3\nvalue = [3, 0]'),
Text(0.85, 0.38888888888888889, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]'),
Text(0.875, 0.5, 'gini = 0.0\nsamples = 120\nvalue = [120, 0]'),
Text(0.95, 0.6111111111111112, 'x[1] <= 247.0\ngini = 0.444\nsamples = 6\nvalue =
[4, 2]'),
Text(0.925, 0.5, 'gini = 0.0\nsamples = 4\nvalue = [4, 0]'),
Text(0.975, 0.5, 'gini = 0.0\nsamples = 2\nvalue = [0, 2]'),
Text(0.95, 0.7222222222222222, 'gini = 0.0\nsamples = 1\nvalue = [0, 1]')

```

Performing Linear Regression using Neural Network

```

## train the algorithm
from sklearn.linear_model import LinearRegression

linreg = LinearRegression()
linreg.fit(X_train, y_train)

# make predictions
y_pred = linreg.predict(X_test)

```

```
# evaluation
from sklearn.metrics import mean_squared_error, r2_score
print('mse=', mean_squared_error(y_test, y_pred))
print('correlation=', r2_score(y_test, y_pred))

mse= 0.08450675765478677
correlation= 0.6563241219440358
```

MSE came out to be 0.08 that is very low with correlation being around 0.66

```
# scale the data using Python and pandas functionality
```

```
mean = X_train.mean(axis=0)
X_train -= mean
std = X_train.std(axis=0)
X_train /= std

X_test -= mean
X_test /= std
```

```
# scale the data using sklearn functionality
from sklearn import preprocessing
```

```
scaler = preprocessing.StandardScaler().fit(X_train)
```

```
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

```
# train the algorithm using 6 then three layers with iteration 500
from sklearn.neural_network import MLPRegressor
```

```
regr = MLPRegressor(hidden_layer_sizes=(6, 3), max_iter=500, random_state=1234)
regr.fit(X_train, y_train)
```

▼ **MLPRegressor**

```
MLPRegressor(hidden_layer_sizes=(6, 3), max_iter=500, random_state=1234)
```

```
# Predicting using neural network
y_pred = regr.predict(X_test)

# evaluation
from sklearn.metrics import mean_squared_error, r2_score
print('mse=', mean_squared_error(y_test, y_pred))
print('correlation=', r2_score(y_test, y_pred))

mse= 0.08174039197427181
correlation= 0.6675745021581087
```

Neural network shows slight improvement after using the scaled data.

Now try using classification

```
from sklearn import preprocessing

scaler = preprocessing.StandardScaler().fit(X_train)

X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)

# train
from sklearn.neural_network import MLPClassifier

clf = MLPClassifier(solver='lbfgs', hidden_layer_sizes=(5, 2), max_iter=500, random_state=1234)
clf.fit(X_train_scaled, y_train)
```

```
▼ MLPClassifier
MLPClassifier(hidden_layer_sizes=(5, 2), max_iter=500, random_state=1234,
              solver='lbfgs')
```

```
# make predictions

pred = clf.predict(X_test_scaled)
```

```
.. . . .
```

```
# output results

print('accuracy = ', accuracy_score(y_test, pred))

confusion_matrix(y_test, pred)
```

```
accuracy = 0.8589743589743589
array([[39,  5],
       [ 6, 28]])
```

```
from sklearn.metrics import classification_report
print(classification_report(y_test, pred))
```

	precision	recall	f1-score	support
0	0.87	0.89	0.88	44
1	0.85	0.82	0.84	34
accuracy			0.86	78
macro avg	0.86	0.85	0.86	78
weighted avg	0.86	0.86	0.86	78

```
# Try a different setting on linear regression that has max iteration of 2500, and hidden
# No significant improvement observed
```

```
regr = MLPRegressor(hidden_layer_sizes=(6, 3), solver='lbfgs', max_iter=2500, random_state=
regr.fit(X_train, y_train)
```

▼

MLPRegressor
MLPRegressor(hidden_layer_sizes=(6, 3), max_iter=2500, random_state=1234, solver='lbfgs')

```
y_pred = regr.predict(X_test)

print('mse=', mean_squared_error(y_test, y_pred))
print('correlation=', r2_score(y_test, y_pred))
```

```
mse= 0.08364533298943039  
correlation= 0.6598274024681188
```

According to the neural network accuracy result, classification has a slight better result on the same data with 0.86 comparing to using linear regression that has correlation of 0.67. That might be because of classification that is most likely because we have more qualitative data after categorical columns and classification was better suited for this specific dataset.

Neural network V.S Regular linear regression Used two settings on neural network regression method to train the Regressor. After comparing to the two models, there are no significant improvements after adjusting max-iteration and different hidden layers. Standard linear regression gives out $mse = 0.08450675765478677$ $correlation = 0.6563241219440358$ and after scaling with neural network. Neural network linear regression model has $mse = 0.08174039197427181$ $correlation = 0.6675745021581087$

In R the data there wasn't as precise of operation on data set as SKLearn using Python but there were some similar features such as transforming data using "Categorical data v.s factors in R. I personally liked using Python better as there were many efficient algorithms importing in SKLearn. As well as much precise matrix multiplication and array methods;